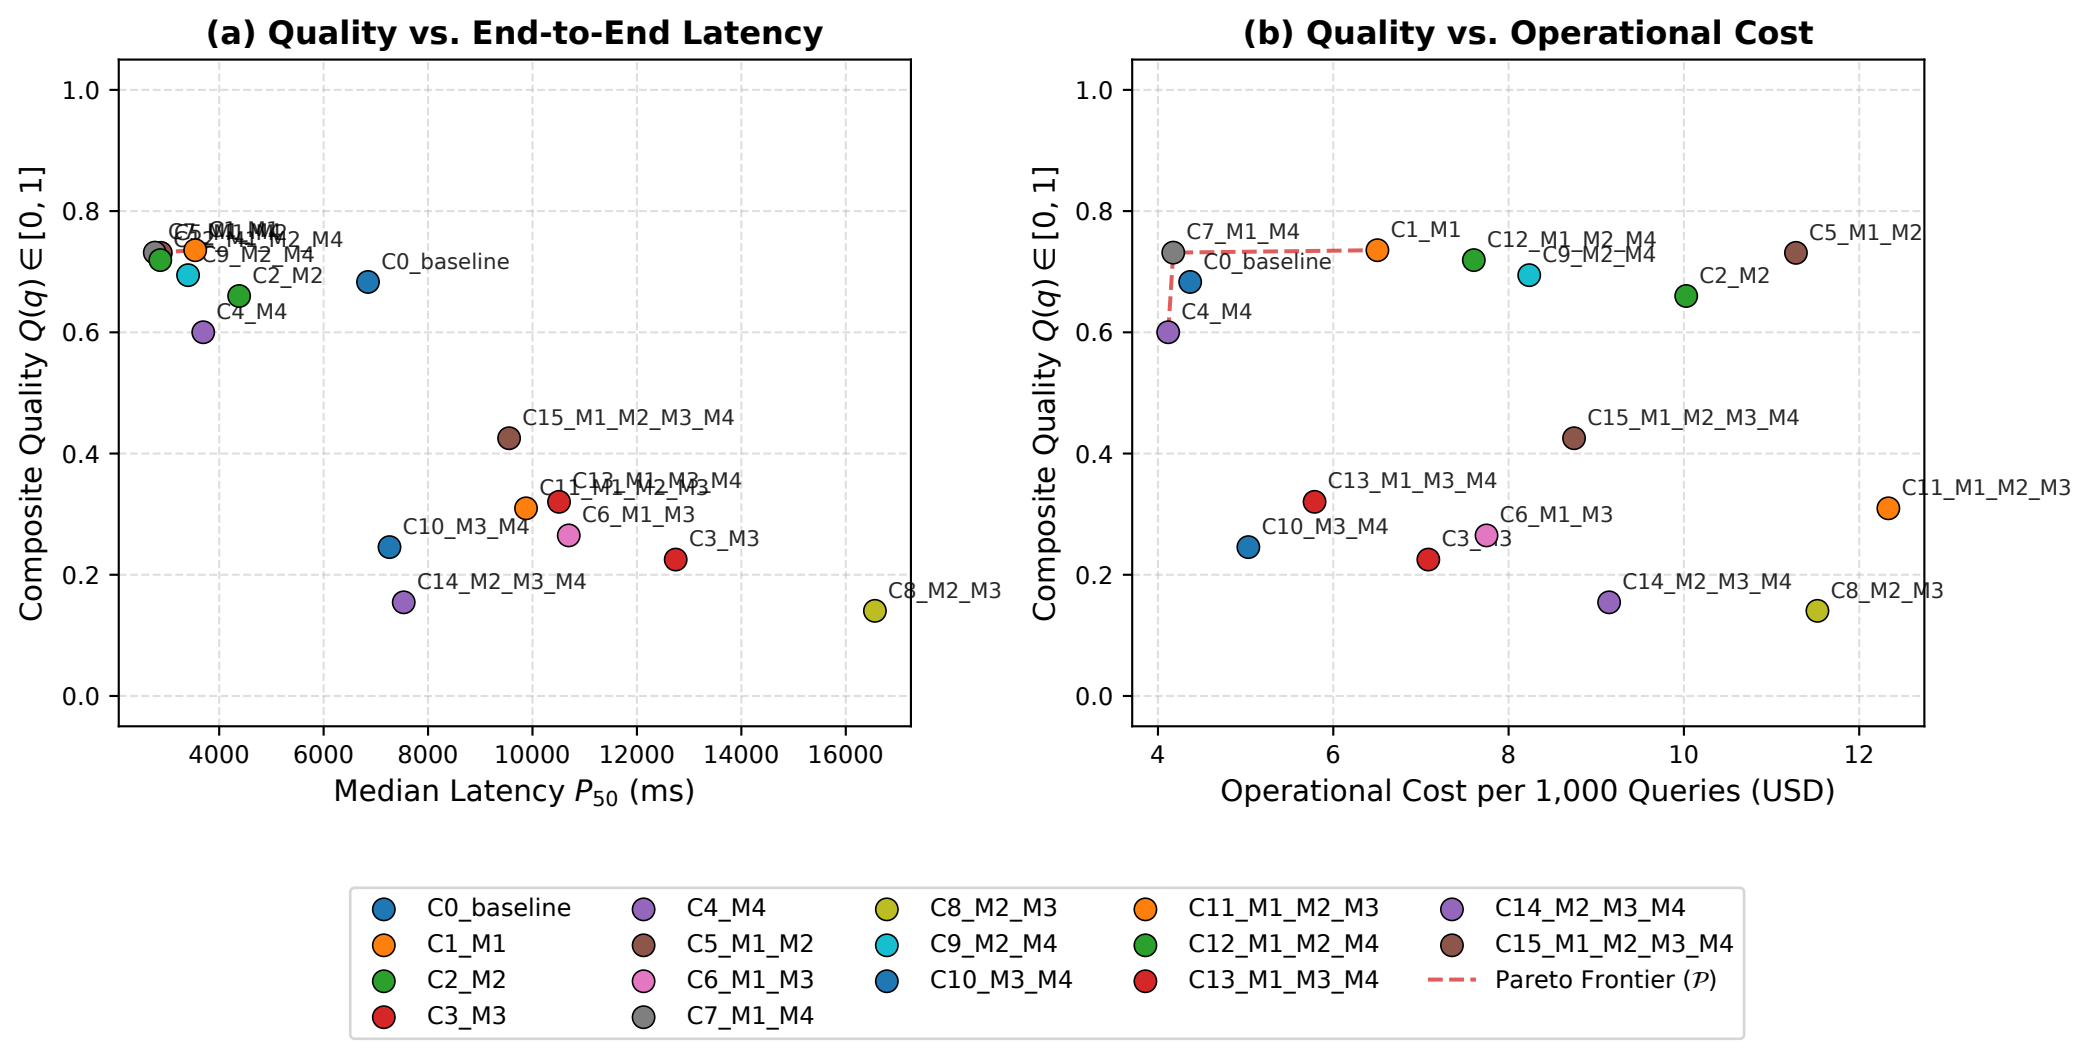
(a) Quality vs. End-to-End Latency
Composite Quality $Q(q) \in [0, 1]$
Median Latency P_{50} (ms)
1.0
0.8
0.6
0.4
0.2
0.0
4000
6000
8000
10000
12000
14000
16000
C0_baseline
C1_M1
C2_M2
C3_M3
C4_M4
C5_M1_M2
C6_M1_M3
C7_M1_M4
C8_M2_M3
C9_M2_M4
C10_M3_M4
C11_M1_M2_M3
C12_M1_M2_M4
C13_M1_M3_M4
C14_M2_M3_M4
C15_M1_M2_M3_M4
(b) Quality vs. Operational Cost
Composite Quality $Q(q) \in [0, 1]$
Operational Cost per 1,000 Queries (USD)
4
6
8
10
12
Pareto Frontier ($\mathcal{P}$)